

The data scientist's toolbox

James Scott

ECO 395M: Data Mining and Statistical Learning

What is this course about?

This class is about gaining knowledge from raw data. You'll learn to use large and complicated data sets to make better decisions.

A mix of practice and principles:

- ▶ Solid understanding of essential statistical ideas
- ▶ Concrete data-crunching skills
- ▶ Best-practice guidelines.

We'll learn what to trust, how to use it, and how to learn more.

Supervised learning

Think “fancy regression”:

- ▶ Given past data on outcomes y paired with features x , can we find patterns that allow us to predict y using x ?
- ▶ Key characteristic: there is a single privileged outcome y .
- ▶ Example: a house has 3 bedrooms (x_1), 2 bathrooms (x_2), 2100 square feet (x_2), and is located in Hyde Park (x_3). What price (y) should it sell for?

In real life, there might be hundreds or thousands of features. OLS might be where you start, but it's rarely where you stop.

Unsupervised learning

This takes us beyond regression, or anything like it:

- ▶ We still have multivariate data and want to find patterns.
- ▶ But there is no single privileged outcome. (“Everything is y.”)

Examples:

- ▶ “Here's data on the shopping basket of every Whole Foods customer at 6th and Lamar last month. Find some patterns that we can use to improve product placement.”
- ▶ “Here's data on every tweet from all 145K followers of @VitaminWater. Tell us about the different kinds of people who follow this brand.”

An alphabet soup of labels. . .

Statistical learning, data mining, data science, ML, AI . . . there are many labels for what we're doing!

- ▶ Econometrics, statistics: focused on understanding the underlying phenomena and formally quantifying uncertainty.
- ▶ Business analytics, data science, data mining: traditionally focused on pragmatic data-analysis tools for applied prediction problems.
- ▶ Machine learning, pattern recognition, artificial intelligence: focused on algorithms with engineering-style performance guarantees.

About “data mining”...

Among economists, “data mining” is a dirty word.

Example: the “Lucas critique”:

- ▶ Fort Knox has never been robbed.
- ▶ Thus historically, there's a zero correlation between security spending at Fort Knox (x) and the likelihood of being robbed (y).
- ▶ Naive “data mining” conclusion: leave Fort Knox unguarded!

This is a total caricature. We'll strive to give data mining a better reputation :-)

About “data mining”...

Some data-mining tools are familiar, or familiar with a twist:

- ▶ data visualization
- ▶ linear regression
- ▶ p-values
- ▶ bootstrapping

Some are totally new:

- ▶ PCA
- ▶ K-means
- ▶ Tree-based models
- ▶ Models for networks
- ▶ Neural nets

People use these tools everywhere

- ▶ Mining client information: Who buys your stuff, what do they pay, what do they think of your new product?
- ▶ Online behavior: Who is on what websites, what do they buy, how do/can we predict or nudge behavior?
- ▶ Collaborative filtering: predict preferences from people who do what you do; recommender engines.
- ▶ Text mining: Connect blogs/emails/news to sentiment, beliefs, or intent. Parsing unstructured data, e.g. EMR.
- ▶ Big regression: mining data to predict asset prices; using unstructured data as controls in observational studies.

The four pillars of data science

1. Data collection
2. Data cleaning and wrangling
3. Data analysis
4. Summary (figures, prose, etc.)

This class focuses mainly on 3-4, a little on 2, and not at all on 1.

Our core principles

You will analyze a lot of data in this course. Our watchwords are *transparency* and *reproducibility*.

- ▶ The end product: you will write a report with beautiful figures, and someone else will marvel at it.
- ▶ Data science is hard enough already: there is zero room for ambiguity or confusion about data or methods.
- ▶ Any competent person should be able to read your description and reproduce exactly what you did.

Our core principles

You will analyze a lot of data in this course. Our watchwords are *transparency* and *reproducibility*.

- ▶ The end product: you will write a report with beautiful figures, and someone else will marvel at it.
- ▶ Data science is hard enough already: there is zero room for ambiguity or confusion about data or methods.
- ▶ Any competent person should be able to read your description and reproduce exactly what you did.

The ideal: “hit-enter” reproducibility.

- ▶ Someone hits enter; your analyses and figures are reproduced from scratch and merged with prose, before their eyes.
- ▶ We rely on few easily mastered tools to put this ideal into practice: RMarkdown, Python/Jupyter, and git (NB: this is a “bilingual” course)

What is RMarkdown?

RMarkdown is an authoring format that allows you to create dynamic documents with R code embedded in them.

- ▶ Integrate R code, output, and visualizations seamlessly.
- ▶ Generate reports in various formats (HTML, PDF, Word, etc.).
- ▶ Reproducible data science: Share your code and results in one document.
- ▶ Easy to learn and use for data scientists and researchers.
- ▶ Supports interactive elements like Shiny apps and HTML widgets.

These slides were written in RMarkdown. You can find the raw .Rmd file on the website.

Basic RMarkdown workflow

1. Install RStudio (if you haven't already).
2. Create a new RMarkdown document in RStudio.
3. Write your narrative text using Markdown syntax.
4. Embed R code chunks to execute code.
5. “Knit” (i.e. compile) the document to generate the output in the desired format.

Follow the tutorials linked on the class website. Note: you're welcome to use Quarto instead!

Jupyter notebooks

Jupyter Notebooks are interactive computing environments that enable you to create and share documents containing live code, equations, visualizations, and narrative text.

- ▶ Code execution in real-time, providing instant feedback.
- ▶ Mix code, Markdown, LaTeX, and HTML in a single document.
- ▶ Rich media support: Display images, videos, and interactive plots.
- ▶ Facilitates data cleaning, analysis, visualization, and machine learning.
- ▶ Great for teaching, research, data exploration, and data communication.

Summary

- ▶ RMarkdown and Jupyter Notebooks are powerful tools for interactive data analysis and reporting.
- ▶ RMarkdown is (obviously) centered around R and is ideal for generating dynamic reports with R code.
- ▶ Jupyter Notebooks support multiple programming languages, with Python being the most popular.
- ▶ Both tools encourage reproducible research and facilitate collaboration among data scientists and researchers.

You can use either in this course. You'll see me use both!

You can also use Quarto, the new kid on the block in this space.

Git and GitHub

git:

- ▶ software for version control.
- ▶ ideal for collaborative work.
- ▶ the basic unit in the git universe is a *repository*, aka “repo”: a collection of files/directories all related to a single task, project, or piece of software.
- ▶ Example: the class website is a git repo.

Git and GitHub

GitHub:

- ▶ a git repository hosting service.
- ▶ a location to store your code in the cloud and easily sync it across multiple machines and multiple collaborators
- ▶ the coolest place on the Internet :-)

Git and GitHub

The git repo for our class website is stored both on GitHub (the remote copy) and my own computer (the local copy).

Basic workflow:

- ▶ Make changes to files in the local copy of the repo.
- ▶ `commit` those changes, thereby creating a snapshot of the repo at a single moment in time that can always be restored.
- ▶ `push` those changes to remote

Git and GitHub

You can use git either through:

- ▶ the command line in a Unix/Linux shell (the hard-core coder's approach)
- ▶ a graphical front-end (e.g. GitHub Desktop, SourceTree). I strongly recommend this for git first-timers!

You'll turn in your homework by posting a link to your Github repo (either RMarkdown or Jupyter)